

OJAS GRAMOPADHYE

og2186@nyu.edu +1(929)753-6104 github.com/ojasgr1706 linkedin.com/in/ojas-gramopadhye

EDUCATION

New York University, New York

Aug 2024 - May 2026

Master of Science, Computer Science

GPA : 3.93/4

Courses : Artificial Intelligence, Deep Learning, Robot Perception, Big Data, Algorithms, ML Systems & Operations

Indian Institute of Technology Bombay, India

May 2019 - Jul 2023

Bachelor of Technology, Computer Science and Engineering

GPA : 7.74/10

Courses : AI & ML, Speech Recognition, Image Processing, Database Systems, Data Structures, Software Systems

KEY PROJECTS

Multi-subject alignment : Neural Speech Representation Learning

Sept 2025 - Ongoing

Graduate Research | NYU Video Lab

New York University

- Developed a multi-subject neural speech decoding framework that learns subject-invariant ECoG representations via contrastive alignment and adapter-based modeling to improve cross-subject generalization.
- Implemented wav2vec2-style self-supervised pretraining for neural signals, using masked encoder modeling, Gumbel quantization, and Transformer context modeling to enhance downstream phoneme classification.

Erica : AI Tutor agent with GraphRAG

Sept 2025 - Dec 2025

Course Project | Artificial Intelligence

New York University

- Developed a Dockerized LLM tutor agent that ingests multi-format course content, enriches it with LLM metadata, stores corpus artifacts in MongoDB, and pipelines them into GraphRAG for grounded, concept-aware retrieval
- Delivered structured, prerequisite-aware explanations by integrating semantic search, a typed knowledge graph, and Streamlit UI with subgraph visualizations, enabling the agent to answer queries with references

Learning Tool Use Policies with Reinforcement Learning

Sept 2025 - Dec 2025

Course Project | Deep Learning

New York University

- Trained a language-conditioned PPO agent to learn sequential tool-use policies over a discrete action space (Retrieve, Calculate, Summarize, Finish), formulating tool invocation as a MDP and achieving 98% accuracy on unseen entities
- Engineered a paraphrase-robust training pipeline with llm based data augmentation and step-cost rewards, leading the agent to learn optimal multi-step reasoning traces and eliminate inefficient tool usage during inference

WORK EXPERIENCE

Detection and Mitigation of Hallucination in Generative Models | EMNLP '24

May 2023 - May 2024

Research Assistant

IBM Research | IIT Bombay

- Proposed a forward-backward clinical reasoning framework (CLINICR) for medical QA, combining differential diagnosis with autoregressive LLM reasoning to reduce hallucinations over eliminative CoT baselines.
- Fine-tuned LLaMA-2 (7B & 70B) with reward-based verification on a new Subjective Medical QA benchmark, achieving 90% expert-validated diagnostic accuracy. (EMNLP Findings 2024) arXiv: arxiv.org/abs/2403.04890

Japanese OCR Pipeline

May 2022 - Jul 2022

Data Science Intern

Daikin

- Delivered a full OCR pipeline for 3,000+ Japanese documents by building document-cleaning and image-normalization scripts, achieving 0.91 character-level accuracy with a MobileNetV3 + Differentiable Binarization detector
- Wrote preprocessing and evaluation pipeline to streamline dataset preparation and measure OCR performance on proprietary scanned documents

Locust Occurrence Modeling and Prediction

May 2021 - Jul 2021

Data Science Associate | R&D Team

Dtime.ai

- Supported an early-warning system for East African farmers by processing 40+ years geospatial vegetation and weather data from East African region into structured 10x10 square-km grids and handling class imbalance with SMOTE
- Improved predicted locust attack precision by 10% over baselines by performing pixel-level classification and modeling spatiotemporal patterns with a ConvLSTM U-Net model

TECHNICAL SKILLS

Machine Learning: LLMs, RAG systems, AI agents, NLP, deep learning, contrastive learning, representation learning, LoRA fine-tuning

Frameworks & Libraries: PyTorch, TensorFlow, HuggingFace, PyTorch Lightning, LangGraph, LangChain, SentenceTransformers, scikit-learn

Data & Systems: Apache Kafka, Apache Spark, Dask, MongoDB, PostgreSQL, ChromaDB, MLflow, Docker, Linux, Bash

Software Engineering: Python, C++, JavaScript, TypeScript, FastAPI, React, Streamlit, d3.js, Git, REST APIs